

homework3

January 30, 2024

```
[ ]: import pandas
import numpy
import utils
import seaborn
import networkx as nx
import matplotlib.pyplot as plt
import gurobipy as gp
from gurobipy import GRB
from itertools import product, combinations
!lscpu | head -n 17
```

Architecture:	x86_64
CPU op-mode(s):	32-bit, 64-bit
Address sizes:	39 bits physical, 48 bits virtual
Byte Order:	Little Endian
CPU(s):	8
On-line CPU(s) list:	0-7
Vendor ID:	GenuineIntel
Model name:	11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz
CPU family:	6
Model:	140
Thread(s) per core:	2
Core(s) per socket:	4
Socket(s):	1
Stepping:	1
CPU(s) scaling MHz:	81%
CPU max MHz:	4800.0000
CPU min MHz:	400.0000

1 Excercise 1

1.0.1 Linear program

We say that two cities are adjacent if they are in range D of each other.

A city is covered if at least one city with a store is adjacent to it.

The objective is to find minimal placement of stores that covers every city.

```
[ ]: def solve(adjacency_matrix, relaxed=False, verbose=False):
    vtype = GRB.CONTINUOUS if relaxed else GRB.BINARY

    model = gp.Model()
    cities = model.addVars(len(adjacency_matrix), vtype=vtype, lb=0.0, ub=1.
↪0)

    model.setObjective(cities.sum(), GRB.MINIMIZE)

    for a in adjacency_matrix:
        model.addConstr(cities.prod(gp.tuplelist(a)) >= 1)

    model.Params.OutputFlag = verbose
    model.optimize()
    return model
```

1.0.2 Solving integer formulation

```
[ ]: coordinates, distances, cities = utils.load_dataset(frac=0.5, seed=123).values()
len(cities)
```

```
[ ]: 1459
```

```
[ ]: %time model = solve(distances < 50, verbose=True)
```

Gurobi Optimizer version 11.0.0 build v11.0.0rc2 (linux64 - "Fedora Linux 36 (Workstation Edition)")

CPU model: 11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz, instruction set [SSE2|AVX|AVX2|AVX512]

Thread count: 4 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 1459 rows, 1459 columns and 73047 nonzeros

Model fingerprint: 0x9dc2816c

Variable types: 0 continuous, 1459 integer (1459 binary)

Coefficient statistics:

Matrix range [1e+00, 1e+00]

Objective range [1e+00, 1e+00]

Bounds range [1e+00, 1e+00]

RHS range [1e+00, 1e+00]

Found heuristic solution: objective 74.0000000

CPU model: 11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz, instruction set [SSE2|AVX|AVX2|AVX512]

Thread count: 4 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 1459 rows, 1459 columns and 73047 nonzeros

Model fingerprint: 0x9dc2816c

Variable types: 0 continuous, 1459 integer (1459 binary)

Coefficient statistics:

Matrix range [1e+00, 1e+00]

Objective range [1e+00, 1e+00]

Bounds range [1e+00, 1e+00]

RHS range [1e+00, 1e+00]

Found heuristic solution: objective 74.0000000

Presolve removed 489 rows and 398 columns

Presolve time: 0.21s

Presolved: 970 rows, 1061 columns, 37479 nonzeros

Found heuristic solution: objective 69.0000000

Variable types: 0 continuous, 1061 integer (1061 binary)

Root relaxation: objective 4.740519e+01, 2618 iterations, 0.24 seconds (0.33 work units)

Nodes		Current Node			Objective Bounds			Work	
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
	0	0	47.40519	0	376	69.000000	47.40519	31.3%	- 0s
H	0	0				67.00000000	47.40519	29.2%	- 0s
H	0	0				65.00000000	47.40519	27.1%	- 0s
H	0	0				57.00000000	47.40519	16.8%	- 0s
H	0	0				56.00000000	47.40519	15.3%	- 0s
H	0	0				54.00000000	47.40519	12.2%	- 0s
	0	0	47.60304	0	380	54.000000	47.60304	11.8%	- 0s
H	0	0				53.00000000	47.61691	10.2%	- 0s
	0	0	47.61691	0	374	53.000000	47.61691	10.2%	- 0s
H	0	0				52.00000000	47.73772	8.20%	- 1s
	0	0	47.73772	0	379	52.000000	47.73772	8.20%	- 1s
	0	0	47.75281	0	387	52.000000	47.75281	8.17%	- 1s
	0	0	47.86800	0	383	52.000000	47.86800	7.95%	- 1s
	0	0	47.89807	0	384	52.000000	47.89807	7.89%	- 1s
	0	0	47.90371	0	391	52.000000	47.90371	7.88%	- 1s
	0	0	47.90467	0	399	52.000000	47.90467	7.88%	- 1s
	0	0	47.90470	0	397	52.000000	47.90470	7.88%	- 1s
H	0	0				51.00000000	47.90735	6.06%	- 2s
	0	0	47.98472	0	374	51.000000	47.98472	5.91%	- 2s
	0	0	47.99264	0	382	51.000000	47.99264	5.90%	- 2s
	0	0	47.99497	0	374	51.000000	47.99497	5.89%	- 2s
	0	0	47.99569	0	384	51.000000	47.99569	5.89%	- 2s
	0	0	47.99589	0	382	51.000000	47.99589	5.89%	- 2s
	0	0	48.01476	0	407	51.000000	48.01476	5.85%	- 2s
	0	0	48.01838	0	397	51.000000	48.01838	5.85%	- 2s
	0	0	48.01882	0	404	51.000000	48.01882	5.85%	- 2s
	0	0	48.01899	0	406	51.000000	48.01899	5.85%	- 2s
	0	0	48.02812	0	412	51.000000	48.02812	5.83%	- 2s
	0	0	48.03021	0	408	51.000000	48.03021	5.82%	- 2s

	0	0	48.03092	0	406	51.00000	48.03092	5.82%	-	2s
	0	0	48.03096	0	405	51.00000	48.03096	5.82%	-	2s
	0	0	48.03452	0	407	51.00000	48.03452	5.81%	-	3s
	0	0	48.03550	0	408	51.00000	48.03550	5.81%	-	3s
	0	0	48.03578	0	411	51.00000	48.03578	5.81%	-	3s
	0	0	48.04018	0	409	51.00000	48.04018	5.80%	-	3s
	0	0	48.04113	0	411	51.00000	48.04113	5.80%	-	3s
	0	0	48.04145	0	407	51.00000	48.04145	5.80%	-	3s
	0	0	48.04372	0	403	51.00000	48.04372	5.80%	-	3s
	0	0	48.04433	0	418	51.00000	48.04433	5.80%	-	3s
	0	0	48.04448	0	415	51.00000	48.04448	5.80%	-	3s
	0	0	48.04515	0	416	51.00000	48.04515	5.79%	-	3s
	0	0	48.04575	0	416	51.00000	48.04575	5.79%	-	3s
H	0	0				50.0000000	48.04613	3.91%	-	5s
	0	2	48.04613	0	416	50.00000	48.04613	3.91%	-	5s
	301	211	48.49195	5	345	50.00000	48.37386	3.25%	270	10s
	755	386	cutoff	12		50.00000	48.58413	2.83%	229	15s
	1236	397	48.97990	12	314	50.00000	48.68417	2.63%	240	20s
	1606	327	48.98409	11	304	50.00000	48.76623	2.47%	230	25s

Cutting planes:

MIR: 244

Zero half: 1

Mod-K: 2

Explored 2195 nodes (462700 simplex iterations) in 28.28 seconds (63.08 work units)

Thread count was 8 (of 8 available processors)

Solution count 10: 50 51 52 ... 69

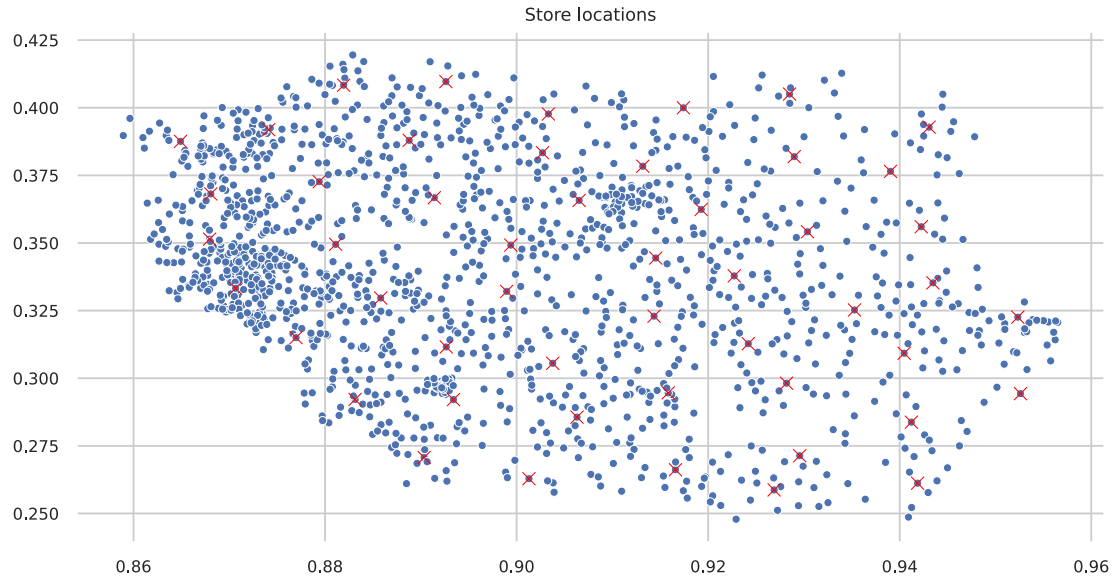
Optimal solution found (tolerance 1.00e-04)

Best objective 5.000000000000e+01, best bound 5.000000000000e+01, gap 0.0000%

CPU times: user 2min 26s, sys: 1.24 s, total: 2min 27s

Wall time: 1min 12s

```
[ ]: stores = coordinates[utils.vars(model, dtype=bool)]
seaborn.scatterplot(x = coordinates[:,0], y = coordinates[:,1], s=20)
seaborn.scatterplot(x = stores[:,0], y = stores[:,1], marker="x", color="red",
    ↪s=50)
plt.title("Store locations"); plt.show()
print(f"There are {model.objVal} stores placed")
```



There are 50.0 stores placed

I was able to compute the solution for at least half the dataset in a reasonable time.
About 50 stores are required to cover all of the cities.

1.0.3 Model relaxation

```
[ ]: %time model = solve(distances < 50, verbose=True, relaxed=True)
```

Gurobi Optimizer version 11.0.0 build v11.0.0rc2 (linux64 - "Fedora Linux 36 (Workstation Edition)")

CPU model: 11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz, instruction set [SSE2|AVX|AVX2|AVX512]

Thread count: 4 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 1459 rows, 1459 columns and 73047 nonzeros

Model fingerprint: 0x3cc0f0be

Coefficient statistics:

Matrix range	[1e+00, 1e+00]
Objective range	[1e+00, 1e+00]
Bounds range	[1e+00, 1e+00]
RHS range	[1e+00, 1e+00]

CPU model: 11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz, instruction set [SSE2|AVX|AVX2|AVX512]

Thread count: 4 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 1459 rows, 1459 columns and 73047 nonzeros

Model fingerprint: 0x3cc0f0be

Coefficient statistics:

Matrix range [1e+00, 1e+00]

Objective range [1e+00, 1e+00]

Bounds range [1e+00, 1e+00]

RHS range [1e+00, 1e+00]

Presolve removed 18 rows and 18 columns

Presolve time: 0.02s

Presolved: 1441 rows, 1441 columns, 71463 nonzeros

Concurrent LP optimizer: dual simplex and barrier

Showing barrier log only...

Ordering time: 0.04s

Barrier statistics:

AA' NZ : 1.024e+05

Factor NZ : 2.611e+05 (roughly 3 MB of memory)

Factor Ops : 5.665e+07 (less than 1 second per iteration)

Threads : 3

Iter	Objective		Residual		Compl	Time
	Primal	Dual	Primal	Dual		
0	1.22398393e+03	0.00000000e+00	0.00e+00	0.00e+00	6.40e-01	0s
1	1.34760819e+02	1.03938869e+01	0.00e+00	6.23e-02	6.69e-02	0s
2	7.71352519e+01	3.19882854e+01	0.00e+00	8.12e-03	1.61e-02	0s
3	5.64557841e+01	4.17874377e+01	0.00e+00	1.14e-03	4.61e-03	0s
4	5.00547000e+01	4.56382592e+01	0.00e+00	6.56e-05	1.39e-03	0s
5	4.80740777e+01	4.67903881e+01	0.00e+00	6.23e-06	4.17e-04	0s
6	4.75743656e+01	4.72646457e+01	0.00e+00	7.77e-16	1.02e-04	0s
7	4.74346072e+01	4.73842751e+01	0.00e+00	5.55e-16	1.68e-05	0s
8	4.74124959e+01	4.73997242e+01	0.00e+00	8.88e-16	4.32e-06	0s
9	4.74066777e+01	4.74036181e+01	0.00e+00	6.66e-16	1.04e-06	0s
10	4.74058637e+01	4.74047201e+01	0.00e+00	4.44e-16	3.88e-07	0s
11	4.74053489e+01	4.74050440e+01	0.00e+00	5.55e-16	1.04e-07	0s
12	4.74052077e+01	4.74051438e+01	0.00e+00	5.55e-16	2.21e-08	0s
13	4.74051957e+01	4.74051830e+01	0.00e+00	5.55e-16	4.40e-09	0s
14	4.74051931e+01	4.74051927e+01	0.00e+00	8.88e-16	1.26e-10	0s

Barrier solved model in 14 iterations and 0.23 seconds (0.15 work units)

Optimal objective 4.74051931e+01

Crossover log...

113 DPushes remaining with DInf 0.0000000e+00 0s

0 DPushes remaining with DInf 0.0000000e+00 0s

154 PPushes remaining with PInf 0.0000000e+00 0s

0 PPushes remaining with PInf 0.0000000e+00 0s

Push phase complete: Pinf 0.0000000e+00, Dinf 9.8674541e-14 0s

Solved with barrier

Iteration	Objective	Primal Inf.	Dual Inf.	Time
238	4.7405193e+01	0.000000e+00	0.000000e+00	0s

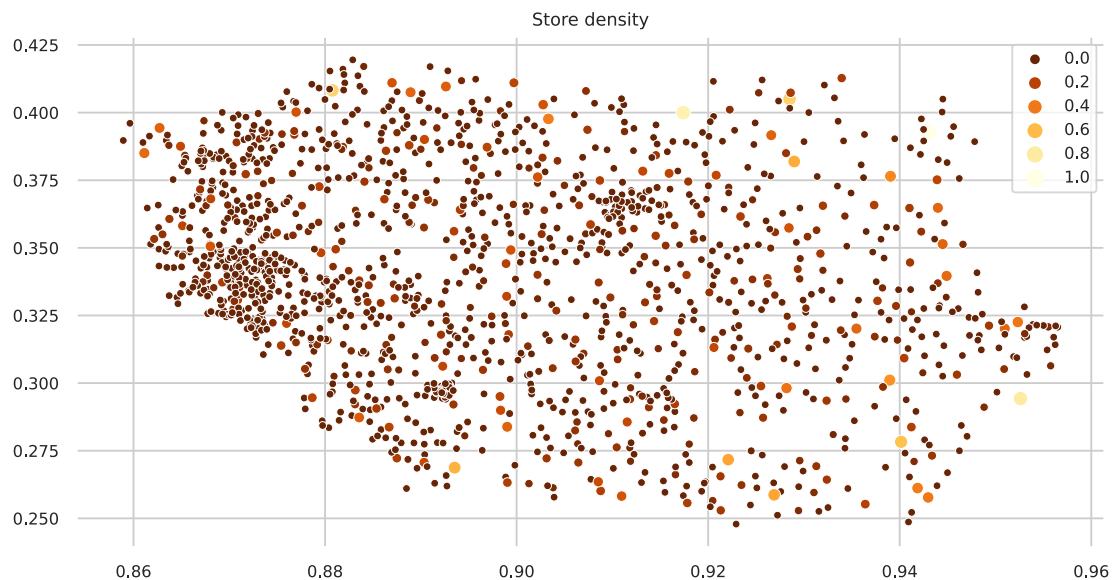
Solved in 238 iterations and 0.30 seconds (0.17 work units)

Optimal objective 4.740519296e+01

CPU times: user 44.3 s, sys: 50.9 ms, total: 44.3 s

Wall time: 44 s

```
[ ]: storiness = utils.vars(model, dtype=float)
seaborn.scatterplot(x = coordinates[:,0], y = coordinates[:,1], s=20,
hue=storiness, size=storiness, palette="YlOrBr_r")
plt.title("Store density"); plt.show()
print(f"There are {model.objVal} stores placed")
```



There are 47.40519295581441 stores placed

The relaxed model is computed in half the time of the original.

In the end we did get away with a little less overall storiness with our continuous formulation, but we are not far off.

1.0.4 Minimal distance per placement size

Unfortunately, I was unable to come up with a much better solution than simply checking different distance ranges since distance is not a variable of my linear program. Such experiment requires the dataset to be much smaller.

```
[ ]: coordinates, distances, cities = utils.load_dataset(frac = 0.05, seed = 123).  
      ↪values()  
      len(cities)
```

```
[ ]: 148
```

```
[ ]: print(  
      solve(distances < 400).ObjVal,  
      solve(distances < 350).ObjVal,  
      solve(distances < 50).ObjVal,  
      )
```

```
1.0 2.0 41.0
```

Objective for $D = 50$ is 41 when considering only 5% of the full dataset.

```
[ ]: pandas.set_option('display.max_columns', None)  
      experiments = [(d, int(solve(distances < d).ObjVal)) for d in range(40, 400)]  
      pandas.DataFrame(experiments, columns=["distance", "k"]).groupby("k").min().  
      ↪reset_index().transpose()
```

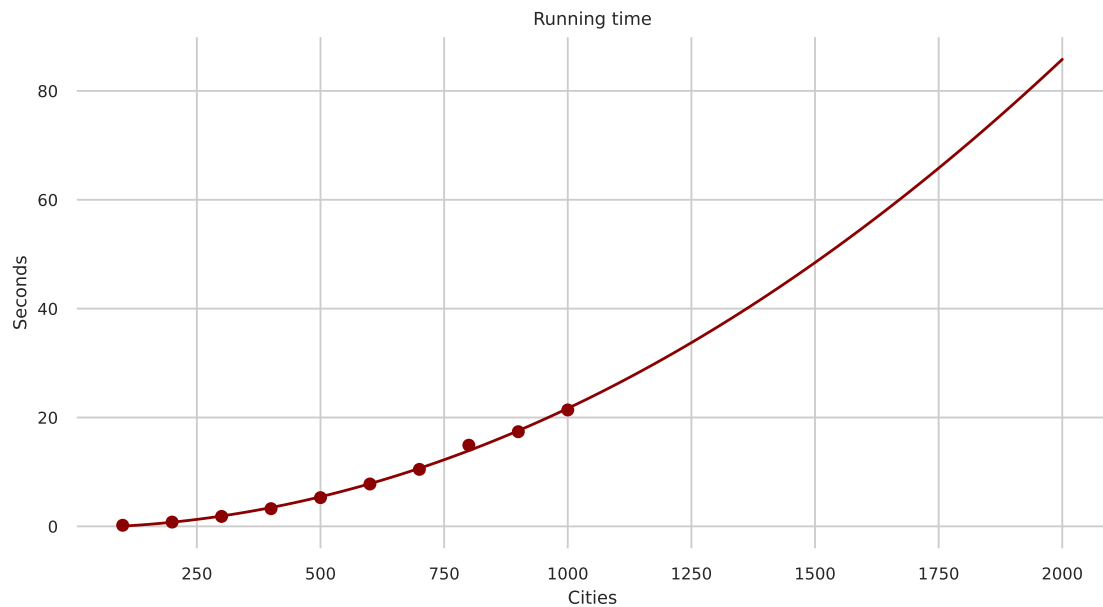
```
[ ]:      0    1    2    3    4    5    6    7    8    9   10   11   12   13  \  
k      1    2    3    4    5    6    7    8    9   10   11   13   14   15  
distance 387  282  226  192  167  147  129  120  115  113  99  96  94  89  
  
      14   15   16   17   18   19   20   21   22   23   24   25   26   27   28   29   30  \  
k      16   17   18   20   22   24   25   26   27   28   30   32   35   36   38   40   41  
distance  87   79   77   73   72   70   67   65   62   60   57   56   55   54   53   52   50  
  
      31   32   33   34   35   36   37   38  
k      42   43   45   47   48   50   53   56  
distance  49   48   46   45   44   43   42   40
```

1.0.5 Running time

```
[ ]: adjacent = tuple(utils.load_dataset(seed=123).values())[1] < 50  
      x = numpy.arange(10) * 100 + 100  
      y = [utils.timeit(lambda: solve(adjacent[:n, :n])) for n in x]  
  
      utils.polyplot(x, y, extrapolate=2.0, degree=2, color="darkred")  
      plt.xlabel('Cities')  
      plt.ylabel('Seconds')  
      plt.title('Running time')
```



```
plt.show()
```



2 Exercise 2

2.0.1 Linear program

Below is an implementation of the [Dantzig-Fulkerson-Johnson](#) linear program formulation for the metric traveling salesman problem.

```
[ ]: def findcycle(order, start = 0):
    cycle = [start]
    while cycle[0] != (next := order[cycle[-1]]):
        cycle.append(next)
    return cycle

def mincycle(order):
    unvisited = set(order)
    cycles = []
    while unvisited:
        start = unvisited.pop()
        cycle = findcycle(order, start)
        cycles.append(cycle)
        unvisited = unvisited.difference(cycle)
    return min(cycles, key=len)

def elimcycles(model, where):
    if where != GRB.Callback.MIPSOL:
```

```

        return

    X = model._vars
    C = utils.cycles(list(model.cbGetSolution(X).values()))

    Q = mincycle(C)
    if len(Q) < len(C):
        model.cbLazy(gp.quicksum(X[edge] for edge in product(Q, Q)) <=
↪len(Q)-1)

def solve(A: numpy.ndarray[int, int], lazy = True, verbose = False):
    model = gp.Model()
    model.Params.OutputFlag = verbose

    V = numpy.arange(n := len(A))
    X = model._vars = model.addVars(*A.shape, vtype=GRB.BINARY)

    model.setObjective( gp.quicksum(X[edge] * A[edge] for edge in
↪product(V, V)) )

    model.addConstrs( X[i, i] == 0 for i in range(n) )

    model.addConstrs( X.sum('*', j) == 1 for j in range(n) )
    model.addConstrs( X.sum(i, '*') == 1 for i in range(n) )

    if lazy:
        model.Params.LazyConstraints = 1
        model.optimize(elimcycles)
    else:
        powerset = utils.powerset(numpy.arange(n))
        model.addConstrs( gp.quicksum(X[edge] for edge in product(Q,
↪Q)) <= len(Q) - 1 for Q in powerset if n != len(Q) >= 2 )
        model.optimize()

    return model

```

2.0.2 Eager implementation

Model has all of the constraints enabled from the beginning.

With this approach I was able to solve instances of at most 16 cities in a reasonable time.

```

[ ]: data = coordinates, distances, cities = utils.load_dataset(frac=0.007,
↪seed=123).values()
len(cities)

```

```

[ ]: 16

```

```
[ ]: %time model = solve(distances, lazy=False, verbose=True)
```

Gurobi Optimizer version 11.0.0 build v11.0.0rc2 (linux64 - "Fedora Linux 36 (Workstation Edition)")

CPU model: 11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz, instruction set [SSE2|AVX|AVX2|AVX512]

Thread count: 4 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 65566 rows, 256 columns and 4456704 nonzeros

CPU model: 11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz, instruction set [SSE2|AVX|AVX2|AVX512]

Thread count: 4 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 65566 rows, 256 columns and 4456704 nonzeros

Model fingerprint: 0x2e06f7f5

Variable types: 0 continuous, 256 integer (256 binary)

Coefficient statistics:

Matrix range [1e+00, 1e+00]

Objective range [3e+01, 7e+02]

Bounds range [1e+00, 1e+00]

RHS range [1e+00, 1e+01]

Found heuristic solution: objective 4467.0000000

Presolve removed 16 rows and 16 columns (presolve time = 5s) ...

Presolve removed 16 rows and 16 columns

Presolve time: 6.35s

Presolved: 65550 rows, 240 columns, 3932400 nonzeros

Variable types: 0 continuous, 240 integer (240 binary)

Root relaxation presolved: 240 rows, 65790 columns, 3932640 nonzeros

Root simplex log...

Iteration	Objective	Primal Inf.	Dual Inf.	Time
0	handle free variables			8s

Starting sifting (using dual simplex for sub-problems)...

Iter	Pivots	Primal Obj	Dual Obj	Time
0	75	-infinity	3.0010000e+03	8s

Sifting complete

118	1.7860000e+03	0.000000e+00	0.000000e+00	9s
118	1.7860000e+03	0.000000e+00	0.000000e+00	9s

Root relaxation: objective 1.786000e+03, 118 iterations, 2.16 seconds (2.26 work

units)

Nodes		Current Node			Objective Bounds			Work	
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
*	0	0		0	1786.0000000	1786.00000	0.00%	-	8s

Explored 1 nodes (118 simplex iterations) in 8.87 seconds (6.30 work units)
Thread count was 8 (of 8 available processors)

Solution count 2: 1786 4467

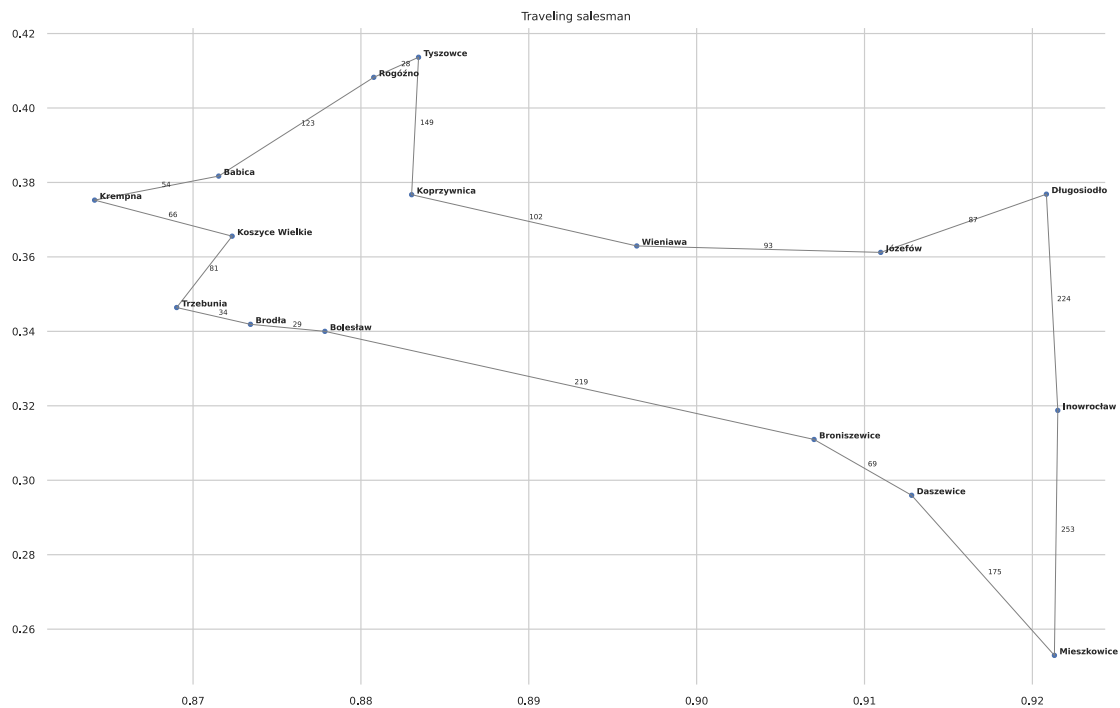
Optimal solution found (tolerance 1.00e-04)

Best objective 1.786000000000e+03, best bound 1.786000000000e+03, gap 0.0000%

CPU times: user 10.4 s, sys: 586 ms, total: 11 s

Wall time: 11.1 s

```
[ ]: utils.graphplot(utils.edges(utils.vars(model)), *data, title='Traveling_↵  
↵salesman')  
print(f"Value of the objective function is {int(model.objVal)}")
```



Value of the objective function is 1786

2.0.3 Lazy implementation

Each time a subcycle is detected we add an extra constraint to the model targeting that specific cycle.

Lazy computation turned out to be quite efficient and thus allowed me to solve instances of up to an order of magnitude greater than before.

```
[ ]: data = coordinates, distances, cities = utils.load_dataset(frac = 0.06, seed = 123).values()  
len(cities)
```

```
[ ]: 172
```

```
[ ]: %time model = solve(distances, lazy=True, verbose=True)
```

Set parameter LazyConstraints to value 1

Gurobi Optimizer version 11.0.0 build v11.0.0rc2 (linux64 - "Fedora Linux 36 (Workstation Edition)")

CPU model: 11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz, instruction set [SSE2|AVX|AVX2|AVX512]

Thread count: 4 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 516 rows, 29584 columns and 59340 nonzeros

Model fingerprint: 0x02487f94

Variable types: 0 continuous, 29584 integer (29584 binary)

Coefficient statistics:

Matrix range [1e+00, 1e+00]

Objective range [3e+00, 7e+02]

Bounds range [1e+00, 1e+00]

RHS range [1e+00, 1e+00]

Gurobi Optimizer version 11.0.0 build v11.0.0rc2 (linux64 - "Fedora Linux 36 (Workstation Edition)")

CPU model: 11th Gen Intel(R) Core(TM) i7-1185G7 @ 3.00GHz, instruction set [SSE2|AVX|AVX2|AVX512]

Thread count: 4 physical cores, 8 logical processors, using up to 8 threads

Optimize a model with 516 rows, 29584 columns and 59340 nonzeros

Model fingerprint: 0x02487f94

Variable types: 0 continuous, 29584 integer (29584 binary)

Coefficient statistics:

Matrix range [1e+00, 1e+00]

Objective range [3e+00, 7e+02]

Bounds range [1e+00, 1e+00]

RHS range [1e+00, 1e+00]

Presolve removed 172 rows and 172 columns

Presolve time: 0.03s

Presolved: 344 rows, 29412 columns, 58824 nonzeros

Variable types: 0 continuous, 29412 integer (29412 binary)

Root relaxation: objective 4.212000e+03, 275 iterations, 0.00 seconds (0.00 work units)

Nodes		Current Node			Objective Bounds			Work		
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time	
0	0	4212.00000	0	-	-	4212.00000	-	-	0s	
0	0	4239.00000	0	-	-	4239.00000	-	-	0s	
0	0	4262.00000	0	-	-	4262.00000	-	-	0s	
0	0	4268.00000	0	-	-	4268.00000	-	-	0s	
0	0	4350.00000	0	38	-	4350.00000	-	-	0s	
0	0	4395.00000	0	38	-	4395.00000	-	-	0s	
0	2	4398.00000	0	38	-	4398.00000	-	-	0s	
1399	1038	5168.00000	41	38	-	4422.00000	-	11.4	5s	
1816	1282	5082.00000	23	-	-	5051.75000	-	12.4	10s	
*	2551	1390	45	5260.0000000	5051.75000	3.96%	13.2	13s		
	2675	1452	5195.75000	29	45	5260.00000	5059.00000	3.82%	13.4	15s
H	2697	1353		5240.0000000	5059.00000	3.45%	13.4	15s		
H	2777	1304		5235.0000000	5060.00000	3.34%	13.5	15s		
	4799	1526	5224.00000	33	83	5235.00000	5134.50000	1.92%	15.5	21s
*	4817	756	38	5207.0000000	5134.50000	1.39%	15.5	21s		
*	5583	805	38	5206.0000000	5164.00000	0.81%	16.2	22s		
*	5589	740	38	5204.0000000	5164.00000	0.77%	16.2	22s		
*	5893	707	34	5202.0000000	5166.00000	0.69%	16.4	23s		
	6922	598	5197.75000	29	226	5202.00000	5178.50000	0.45%	17.1	25s
*	7280	467	31	5199.0000000	5180.00000	0.37%	17.3	25s		

Cutting planes:

Gomory: 8

Cover: 7

Inf proof: 2

Mod-K: 5

Lazy constraints: 103

Explored 8102 nodes (142216 simplex iterations) in 26.70 seconds (28.67 work units)

Thread count was 8 (of 8 available processors)

Solution count 8: 5199 5202 5204 ... 5260

Optimal solution found (tolerance 1.00e-04)

Best objective 5.199000000000e+03, best bound 5.199000000000e+03, gap 0.0000%

User-callback calls 19510, time in user-callback 3.77 sec

CPU times: user 1min 35s, sys: 2.19 s, total: 1min 37s

Wall time: 27.1 s

```
[ ]: utils.graphplot(utils.edges(utils.vars(model)), *data, title='Traveling_
↪salesman')
print(f"Value of the objective function is {int(model.objVal)}")
```



Value of the objective function is 5199

3 Exercise 3

3.0.1 Christofides algorithm

Below implementation uses the [networkx](#) library.

The minimal weight matching implementation is claimed to be in $O(|V|^3)$ time.

```
[ ]: def christofides(distances):
    G = nx.Graph()
    G.add_nodes_from(numpy.arange(len(distances)))
    for edge in combinations(G.nodes, 2):
        G.add_edge(*edge, weight=distances[edge])

    mst = nx.minimum_spanning_tree(G)
    ods = G.subgraph([v for v, degree in mst.degree if degree % 2])

    MG = nx.MultiGraph()
    MG.add_edges_from(mst.edges)
    MG.add_edges_from(nx.min_weight_matching(ods))
```

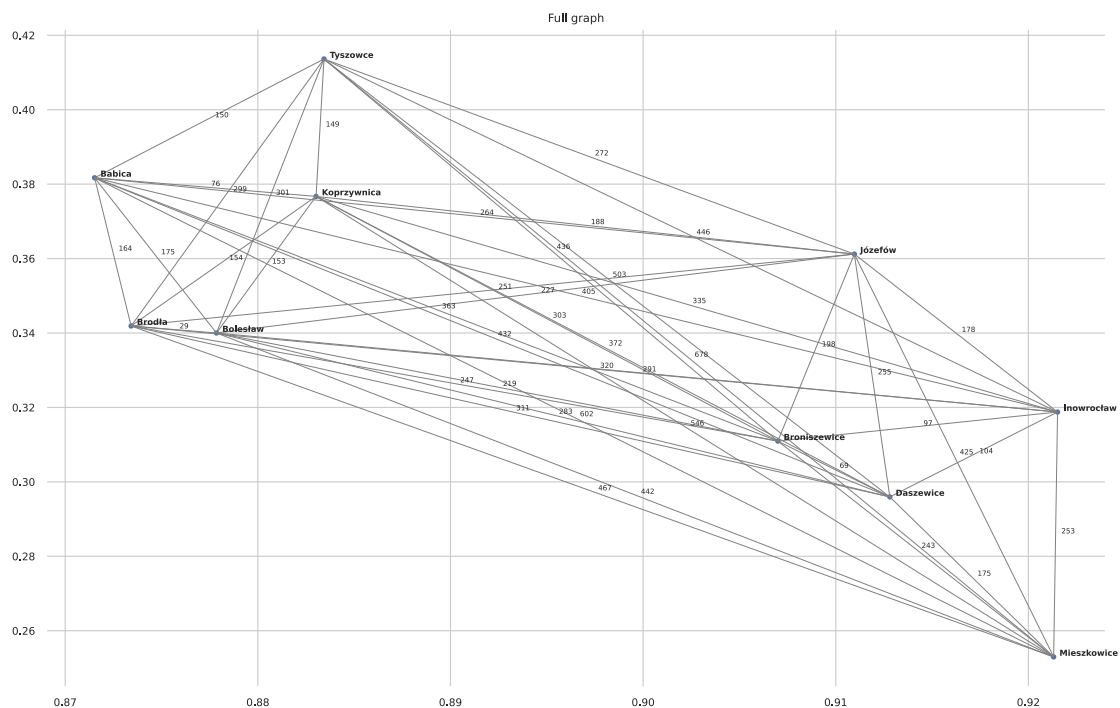
```
return sum(distances[edge] for edge in utils.tour(MG))
```

3.0.2 Christofides step by step

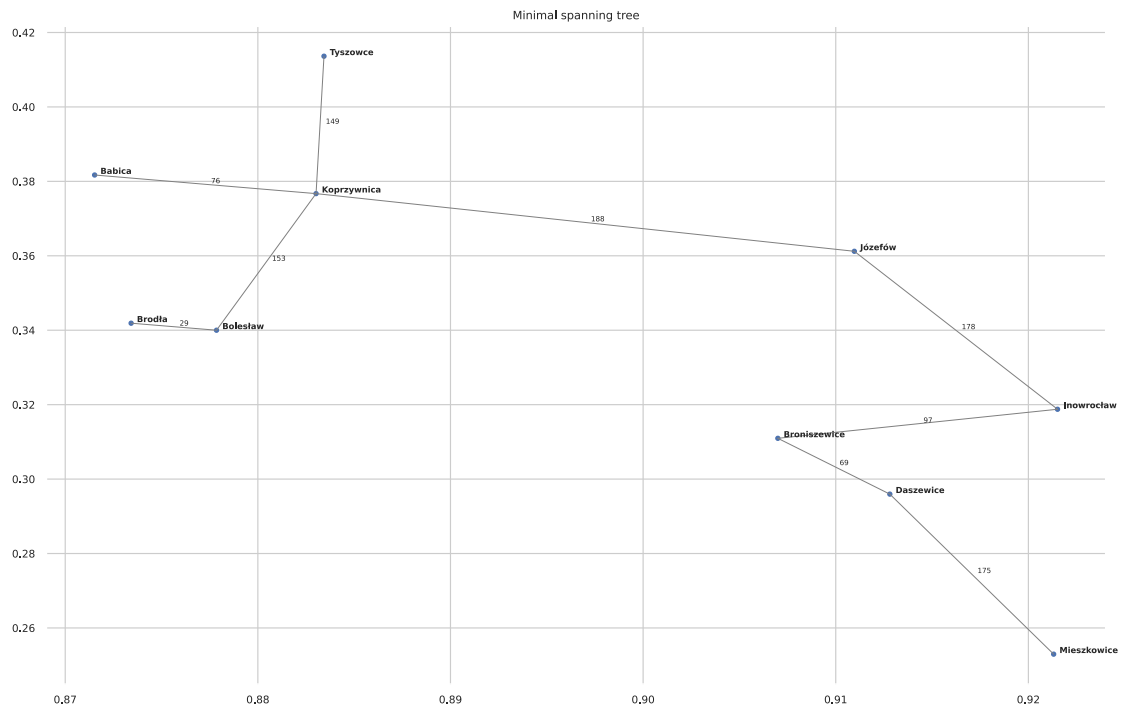
```
[ ]: data = coordinates, distances, cities = utils.load_dataset(frac = 0.005, seed = 123).values()
len(cities)
```

```
[ ]: 10
```

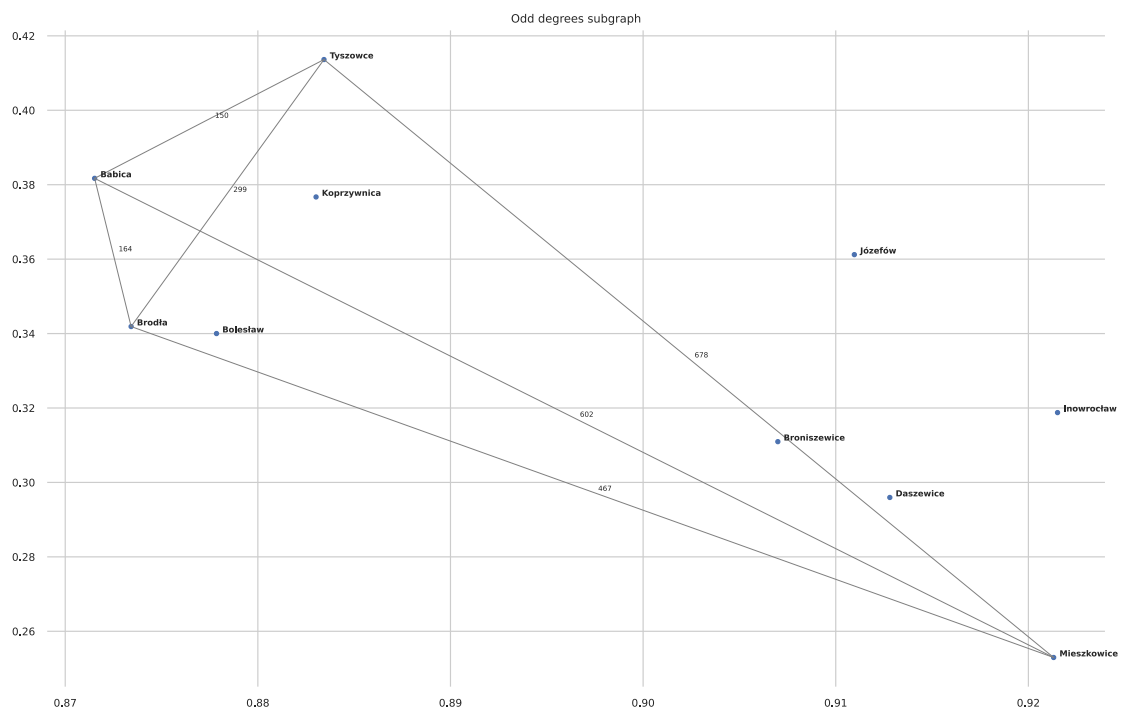
```
[ ]: G = nx.Graph()
G.add_nodes_from(numpy.arange(n := len(cities)))
for edge in combinations(G.nodes, 2):
    G.add_edge(*edge, weight=distances[edge])
utils.graphplot(G.edges, *data, title=f"Full graph")
```



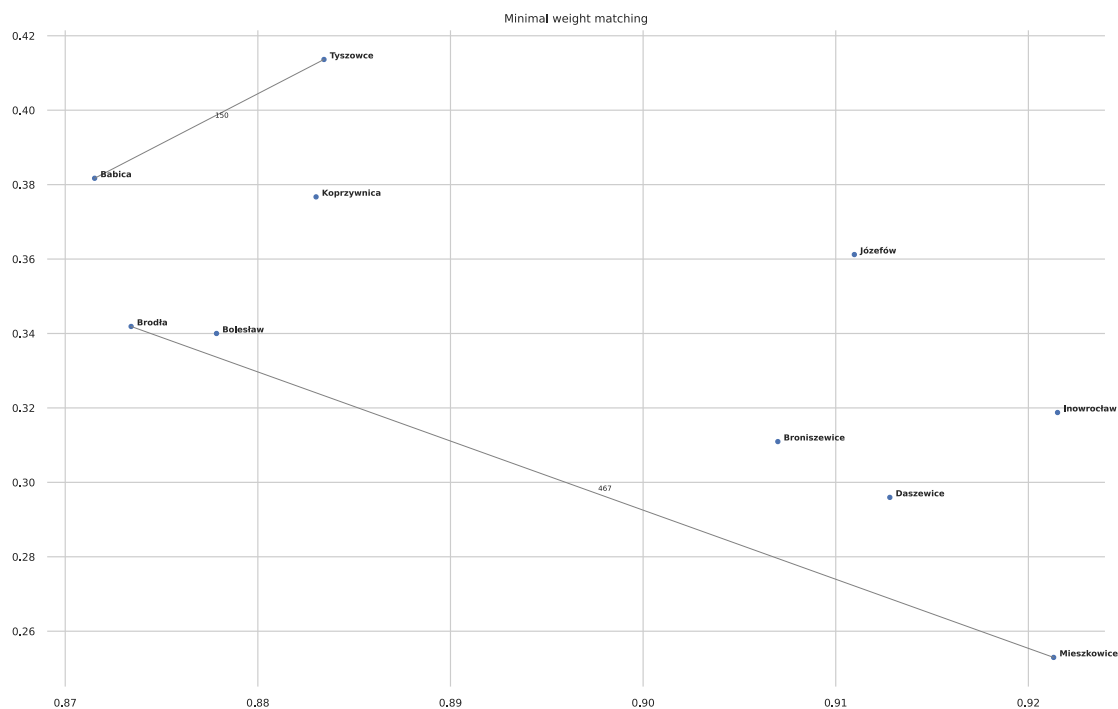
```
[ ]: utils.graphplot((MST := nx.minimum_spanning_tree(G)).edges, *data,
title="Minimal spanning tree")
```

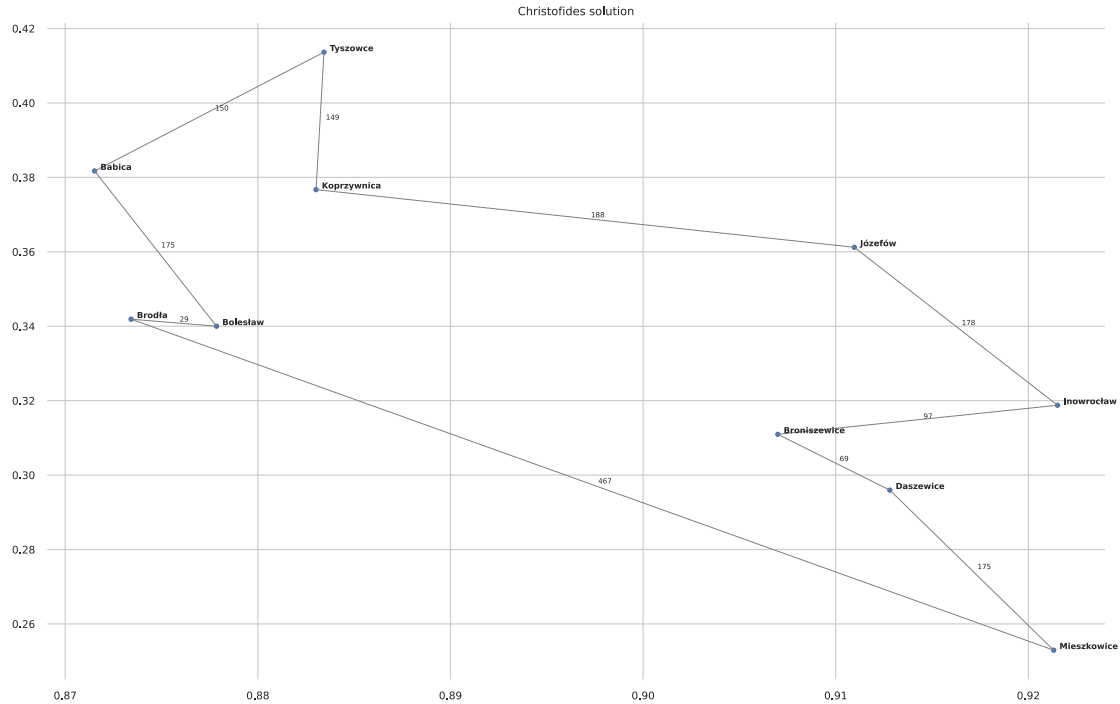
```
[ ]: utils.graphplot((ODS := G.subgraph([v for v, degree in MST.degree if degree % 2 == 1])).edges, *data, title = "Odd degrees subgraph")
```



```
[ ]: utils.graphplot((MWM := ODS.edge_subgraph(nx.min_weight_matching(ODS))).edges,
↳*data, title="Minimal weight matching")
```



```
[ ]: MG = nx.MultiGraph()
MG.add_edges_from(MST.edges)
MG.add_edges_from(MWM.edges)
utils.graphplot(utils.tour(MG), *data, title="Christofides solution")
```



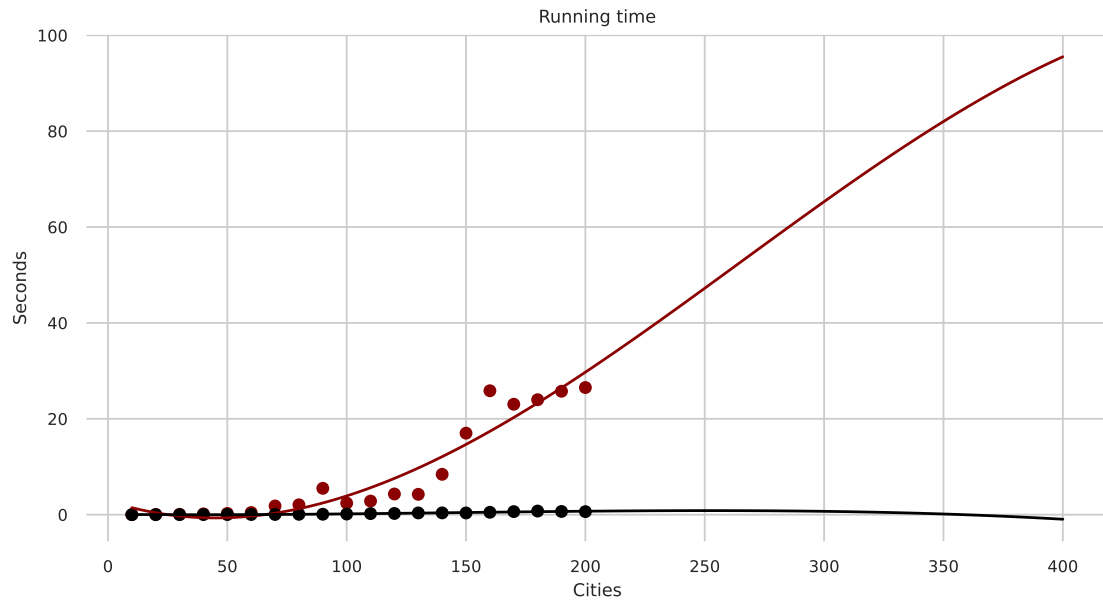
3.0.3 Efficiency and running time

```
[ ]: data = coordinates, distances, cities = utils.load_dataset(frac = 0.06, seed = 123).values()
len(cities)
```

```
[ ]: 172
```

```
[ ]: x = numpy.arange(20) * 10 + 10
y = [utils.timeit(lambda: solve(distances[:n, :n])) for n in x]
z = [utils.timeit(lambda: christofides(distances[:n, :n])) for n in x]

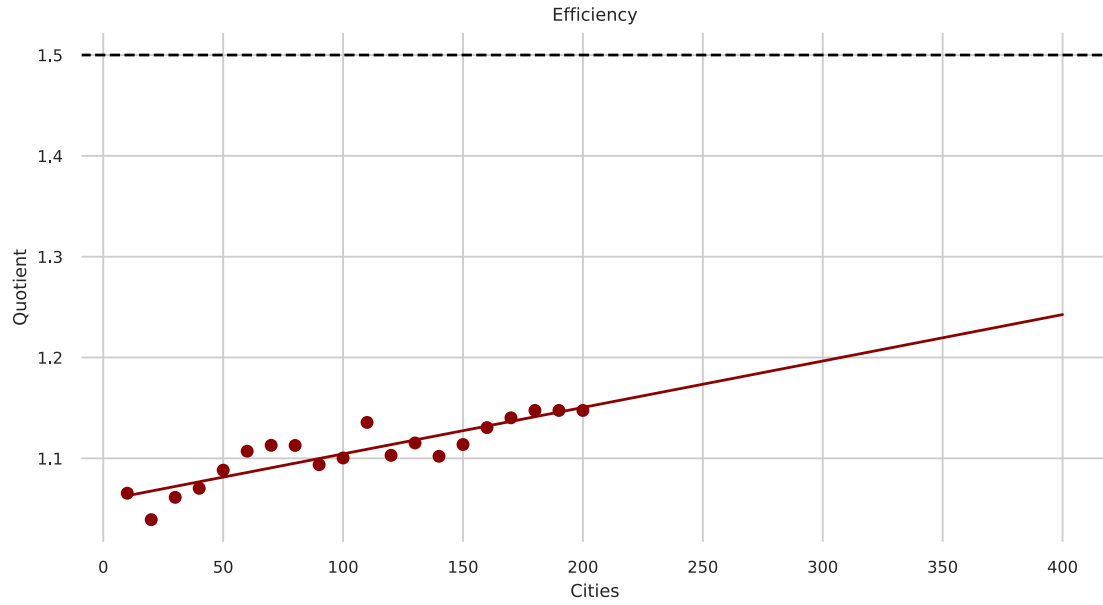
utils.polyplot(x, y, extrapolate=2.0, degree=3, color="darkred")
utils.polyplot(x, z, extrapolate=2.0, degree=3, color="black")
plt.xlabel('Cities')
plt.ylabel('Seconds')
plt.title('Running time')
plt.show()
```



Since our Christofides implementation works in polynomial time while the linear program approach is ultimately exponential this is pretty much the result we would expect.

```
[ ]: x = numpy.arange(20) * 10 + 10
y = [christofides(distances[:n, :n]) / solve(distances[:n, :n]).objVal for n in range(20)]

utils.polyplot(x, y, extrapolate=2.0, degree=1, color="darkred")
plt.axhline(y=3/2, color='black', linestyle='--')
plt.xlabel('Cities')
plt.ylabel('Quotient')
plt.title('Efficiency')
plt.show()
```



We see that the approximation error grows larger with the sample size but by the guarantees of the christofides approximation ratio we expect it to never cross the 1.5 boundary.